

Attention Is All You Need

Ashish Vaswani*
Google Brain
avaswani@google.com

Noam Shazeer*
Google Brain
noam@google.com

Niki Parmar*
Google Research
nikip@google.com

Jakob Uszkoreit*
Google Research
usz@google.com

Llion Jones*
Google Research
llion@google.com

Aidan N. Gomez*[†]
University of Toronto
aidan@cs.toronto.edu

Łukasz Kaiser*
Google Brain
lukaszkaiser@google.com

Illia Polosukhin*[‡]
illia.polosukhin@gmail.com

[arXiv Link to the paper](#)

Published in **Neural Information Processing Systems(NIPS)**, 2017

Presented by **Partha Pratim Saha**

19 Feb 2020

Problems that we want to solve using AI

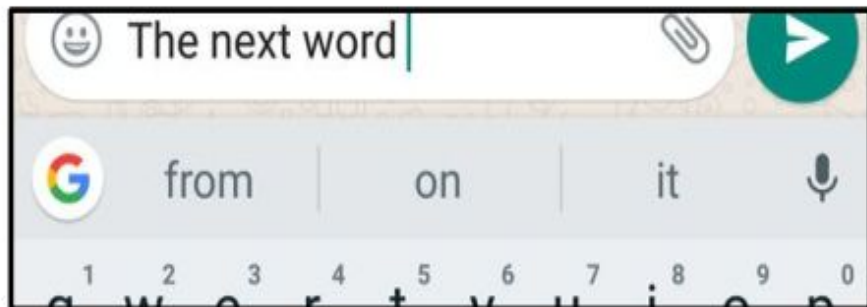
Various real life applications:

- ★ **Machine translation:** Translate a sentence from English to German or French
 - Google translator is using such a model in their production from 2016
- ★ **Language modeling:** Predict next best word when you chat
- ★ **Image Captioning:** Given an image, machine explains about it automatically
- ★ **Advantages:**
 - If your student or client is from another country, AI is helping you in the mentioned scenarios through this model

Examples



1. Group picture of nine tourists and one local on a grey rock with a lake in the background;
2. Five people are standing and four are squatting on a brown rock in the foreground;

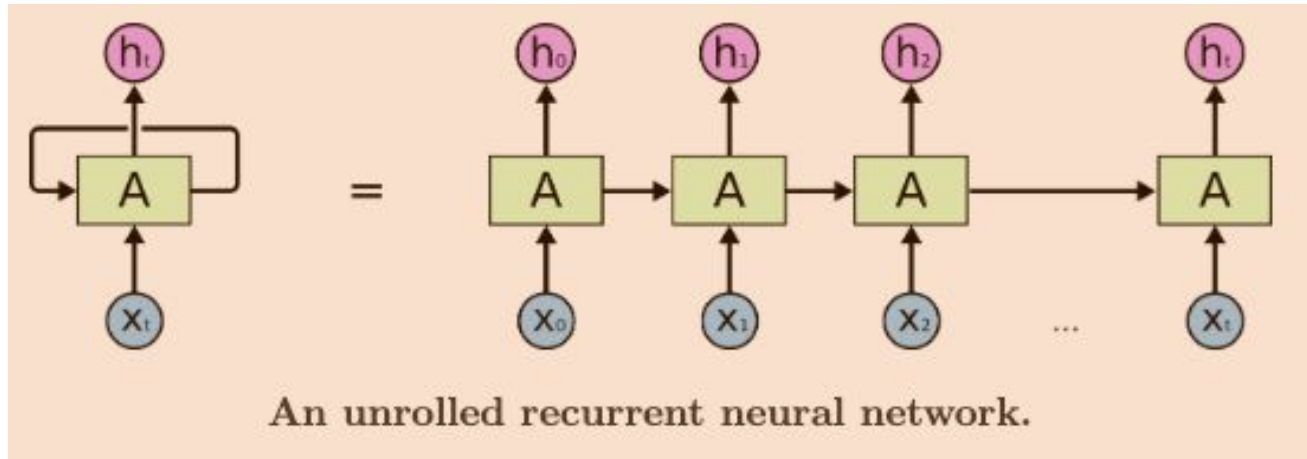


More examples: Our ambiguous language

- "Bob could not put the trophy inside the suitcase because it was too big."
→ What does **it** refers to here: **Trophy OR Suitcase?**
- "Margarett dropped the plate on the table and broke it."
→ What does **it** refers to here: **Table OR Plate?**
- "The animal didn't cross the street because it was too tired."
→ What does **it** refers to here: **Animal OR Street?**
- "The animal didn't cross the street because it was flooded."
→ What does **it** refers to here: **Animal OR Street?**

Existing Deep Learning solution (1)

- **RNN (Recurrent Neural Network)** that is used when the output from previous step is fed as input to the next unit as we see in the below picture:




★ Drawback:

- Vanishing and exploding gradient.
- Prohibits parallelization within instances
- Next step output depends on previous step

Existing Deep Learning solution (2)

- **LSTM (Long short term memory)** is a variant of RNN that is used to remember some part of the previous dependent words contextually.




★ Drawback:

- Cannot parallelization within instances

Existing Deep Learning solution (3)

- **CNN (Convolution Neural Network)** is used for object identification in a given image (spatial data)



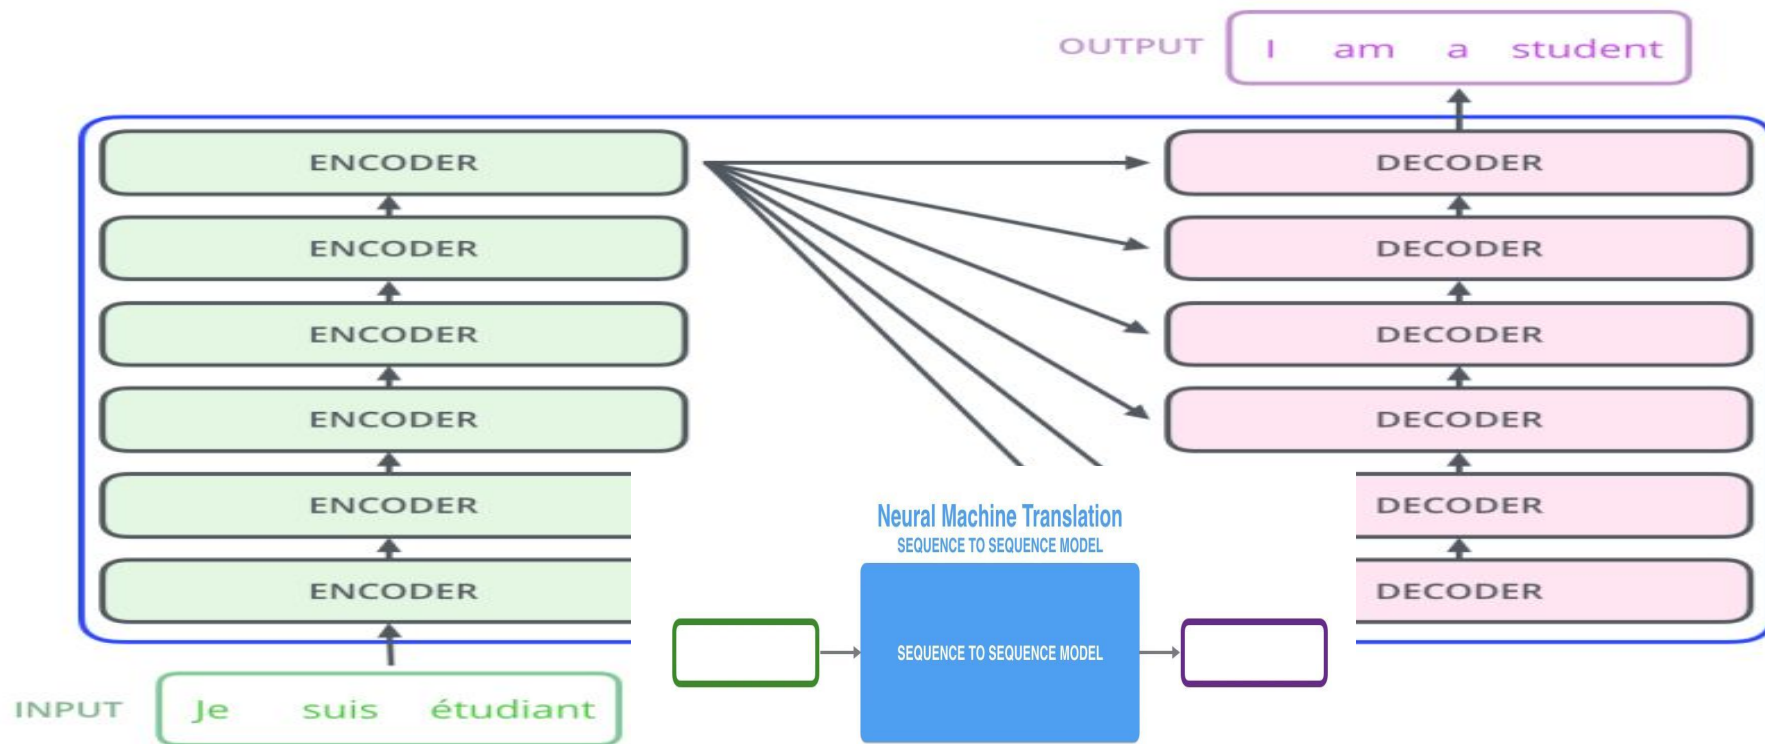
★ Drawback:

- Parallelize within layer but NOT across layers in the network
- Slow running time

Proposed solution: **Transformer** model & some definitions

- **Sequence2Sequence:** A technique that takes a sequence of items like words, letters, features of an images...etc as a vector and outputs another vector as a sequence of items.
- **Encoder:** A function maps an input sequence of symbol representations $(X_1..X_n)$ to a sequence of continuous representations $Z = (Z_1..Z_n)$
- **Decoder:** Given a representation Z , the decoder generates an output sequence $(Y_1..Y_m)$ of symbols one element at a time.
- **Attention:** A function that map a given query, key, value to a probability distribution, where the query, keys, values, and output are vectors.
- **Self Attention:** A method focus on some of the words in the vicinity of the given input sequence
- **Multihead Attention:** Doing self attention mechanism multiple times linearly with different inputs of Q/V/K matrices & have different sets of output matrices to concatenate them together.
- **Output:** Weighted sum of the values, where weight assigned to each value is computed by a function of the query with the corresponding key.

Transformer Model (6 identical encoders and decoders)

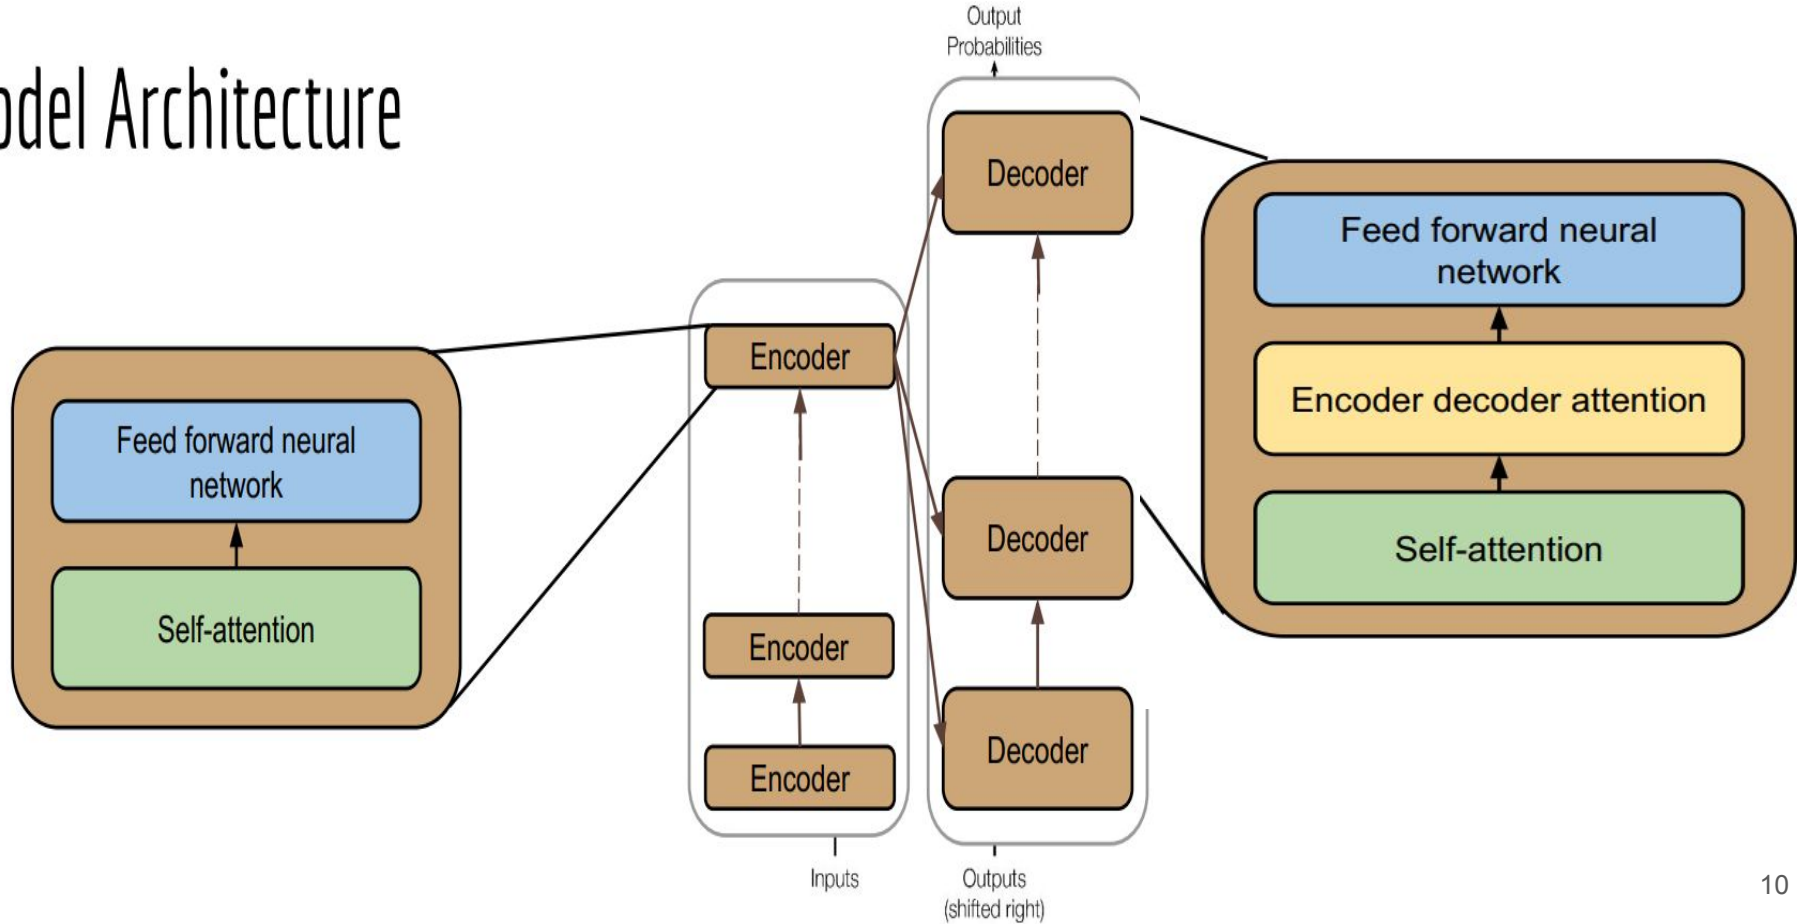


★ Advantages

- Capture long range dependencies using self attention
- Sequentiality parallelize within instances

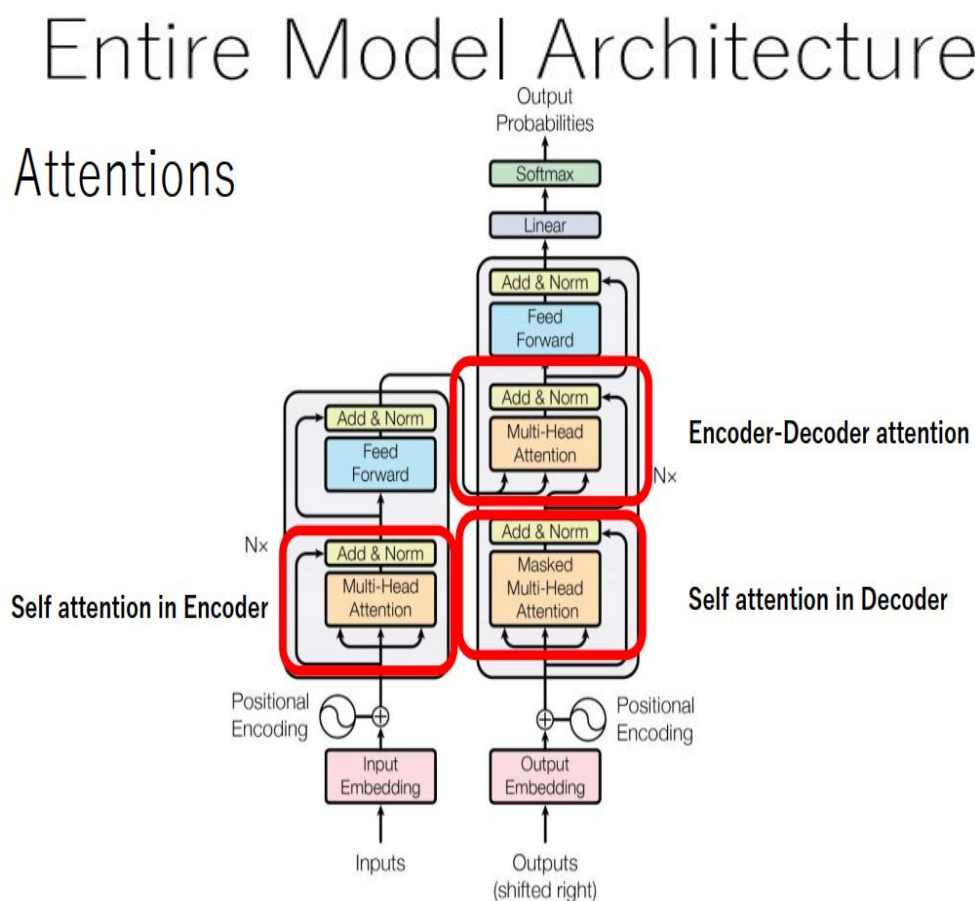
Deep dive into Encoder - Decoder layers

Model Architecture

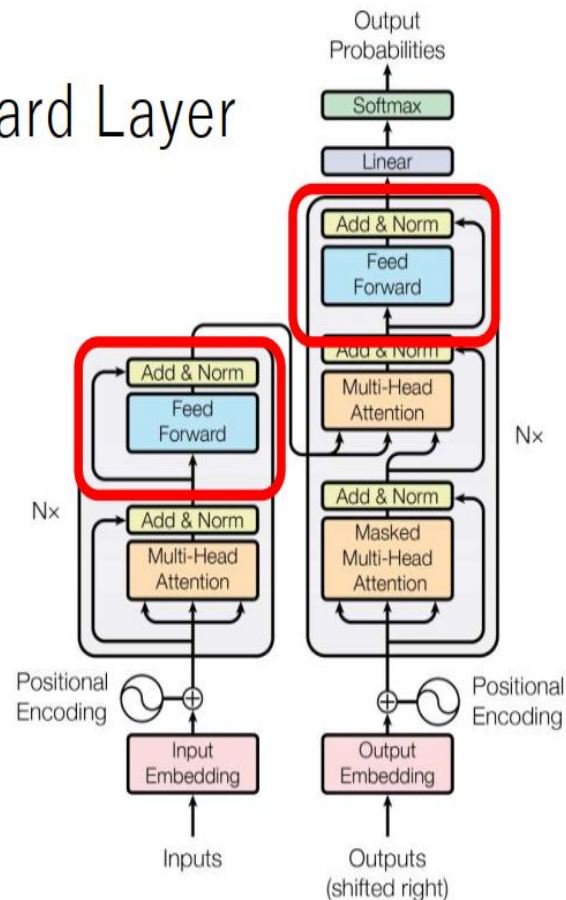


Entire Model Architecture

Attentions



Feed Forward Layer




The **encoder** processes each item in the input sequence, it compiles the information it captures into a vector (called as **context vector**). The top **encoder** sends the **context** to all **decoders**, which begins producing the output sequence item by item.

What is Attention?


Attention : (vector, matrix, matrix) \rightarrow vector

$$\text{Attention}(\text{query}, \text{Key}, \text{Value}) = \text{Softmax}(\text{query} \cdot \text{Key}^T) \cdot \text{Value}$$



The self-attention calculation in matrix form

Scaled Dot-Product Attention



- **Attention** maps a given query, key, value to a probability distribution, where the query, keys, values are 64 dimensional vectors. **Softmax function** is used for multiclass classification as in example work “it” check different words with different probability values. And softmax gives

Why Self Attention

Table 1: Maximum path lengths, per-layer complexity and minimum number of sequential operations for different layer types. n is the sequence length, d is the representation dimension, k is the kernel size of convolutions and r the size of the neighborhood in restricted self-attention.

Layer Type	Complexity per Layer	Sequential Operations	Maximum Path Length
Self-Attention	$O(n^2 \cdot d)$	$O(1)$	$O(1)$
Recurrent	$O(n \cdot d^2)$	$O(n)$	$O(n)$
Convolutional	$O(k \cdot n \cdot d^2)$	$O(1)$	$O(\log_k(n))$
Self-Attention (restricted)	$O(r \cdot n \cdot d)$	$O(1)$	$O(n/r)$

- n = number of words in a longest sentence in the training set ~ 70
- d = number of hidden layers in neural network ~ 1000
- **Sequential operations** = amount of parallel computations
- **Maximum path length** = longest dependency within the sentence

(self attention: $n^2 \cdot d$) $70 \cdot 70 \cdot 1000 \ll 70 \cdot 1000 \cdot 1000$ (RNN or CNN)

Training the model: Experiment

- Data
 - WMT2014 English-German : 4.5 million sentence pairs
 - WMT2014 English-French : 36 million sentences
- Hardware and Schedule
 - 8 NVIDIA P100 GPUs
 - Base model : 100,000 steps or 12 hours
 - Big model : 300,000 steps (3.5 days)
- Optimizer : Adam
 - Warm up, and then decrease learning rate.
- Trained a 4-layer (4 encoders-decoders) 1024 dimensional transformer model on the **Wall Street Journal (WSJ)** portion about 40K training sentences. They have also trained on **BerkleyParser** corpora from with approximately 17M sentences.

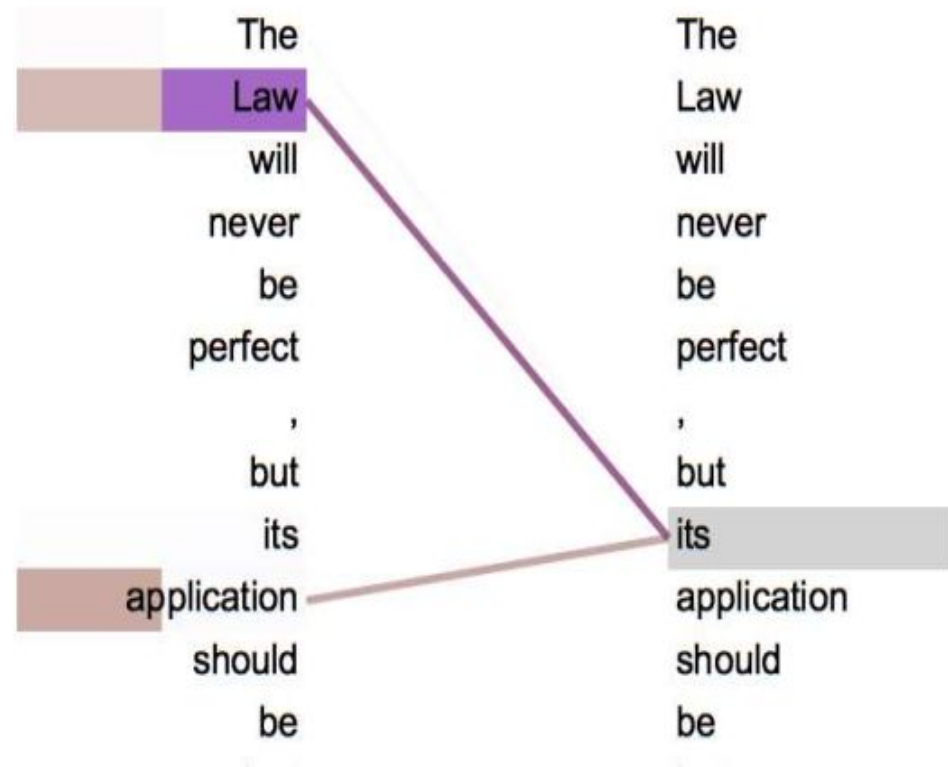
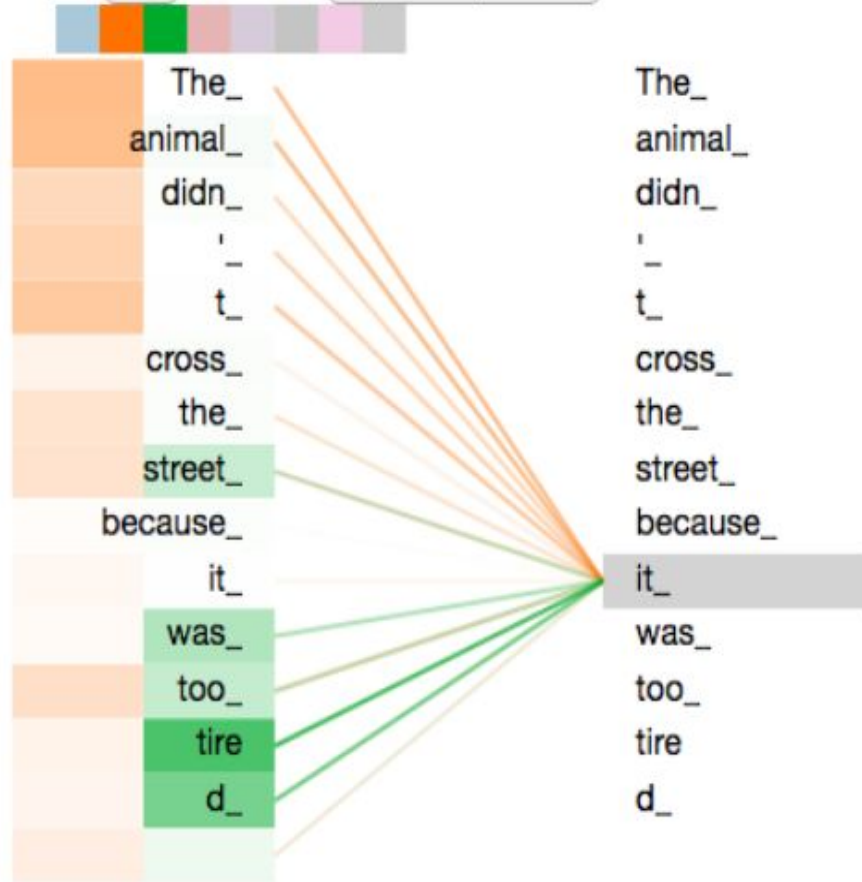
Experiment - Result

Table 2: The Transformer achieves better BLEU scores than previous state-of-the-art models on the English-to-German and English-to-French newstest2014 tests at a fraction of the training cost.

Model	BLEU		Training Cost (FLOPs)	
	EN-DE	EN-FR	EN-DE	EN-FR
ByteNet [15]	23.75			
Deep-Att + PosUnk [32]		39.2		$1.0 \cdot 10^{20}$
GNMT + RL [31]	24.6	39.92	$2.3 \cdot 10^{19}$	$1.4 \cdot 10^{20}$
ConvS2S [8]	25.16	40.46	$9.6 \cdot 10^{18}$	$1.5 \cdot 10^{20}$
MoE [26]	26.03	40.56	$2.0 \cdot 10^{19}$	$1.2 \cdot 10^{20}$
Deep-Att + PosUnk Ensemble [32]		40.4		$8.0 \cdot 10^{20}$
GNMT + RL Ensemble [31]	26.30	41.16	$1.8 \cdot 10^{20}$	$1.1 \cdot 10^{21}$
ConvS2S Ensemble [8]	26.36	41.29	$7.7 \cdot 10^{19}$	$1.2 \cdot 10^{21}$
Transformer (base model)	27.3	38.1	$3.3 \cdot 10^{18}$	
Transformer (big)	28.4	41.0	$2.3 \cdot 10^{19}$	

Self-attention visualization

Layer: 5 ▾ Attention: Input - Input ▾



Summary

- I talked about a novel method “**Transformer**” - the first sequence to sequence machine translation process entirely based on **Attention**, replacing recurrent layers in neural network with **multi-headed self-attention**.
- **Future work** on impactful applications:
 - Current range of **self attention neighbourhood = r (fixed)** in input sequence that is centered around a respective output position. They will look for different values of r for dependencies in a given sentence.
 - To sum up reading comprehension.
 - Abstractive summarization of research papers and news articles.

Thank you for your
Attention

Backup: Methodology (matrix multiplications)

\mathbf{X} = **512** dimensional matrix (vector embedding)

$\mathbf{Q} = \mathbf{K} = \mathbf{V} = \mathbf{dk} = \mathbf{64}$ dimensional matrices



Trainable matrices \mathbf{WQ} , \mathbf{WK} , \mathbf{WV} are used for **8** matrix multiplications in case of **multi-headed attention**

$\mathbf{X1}$ is multiplied by \mathbf{WQ} weight matrix produce $\mathbf{q1}$ **query vector** for that word

$\mathbf{X2}$ is multiplied by \mathbf{WQ} weight matrix produce $\mathbf{q2}$ **query vector** for 2nd word and so on

Matrix $\mathbf{W0}$ is used for **1** matrix multiplication

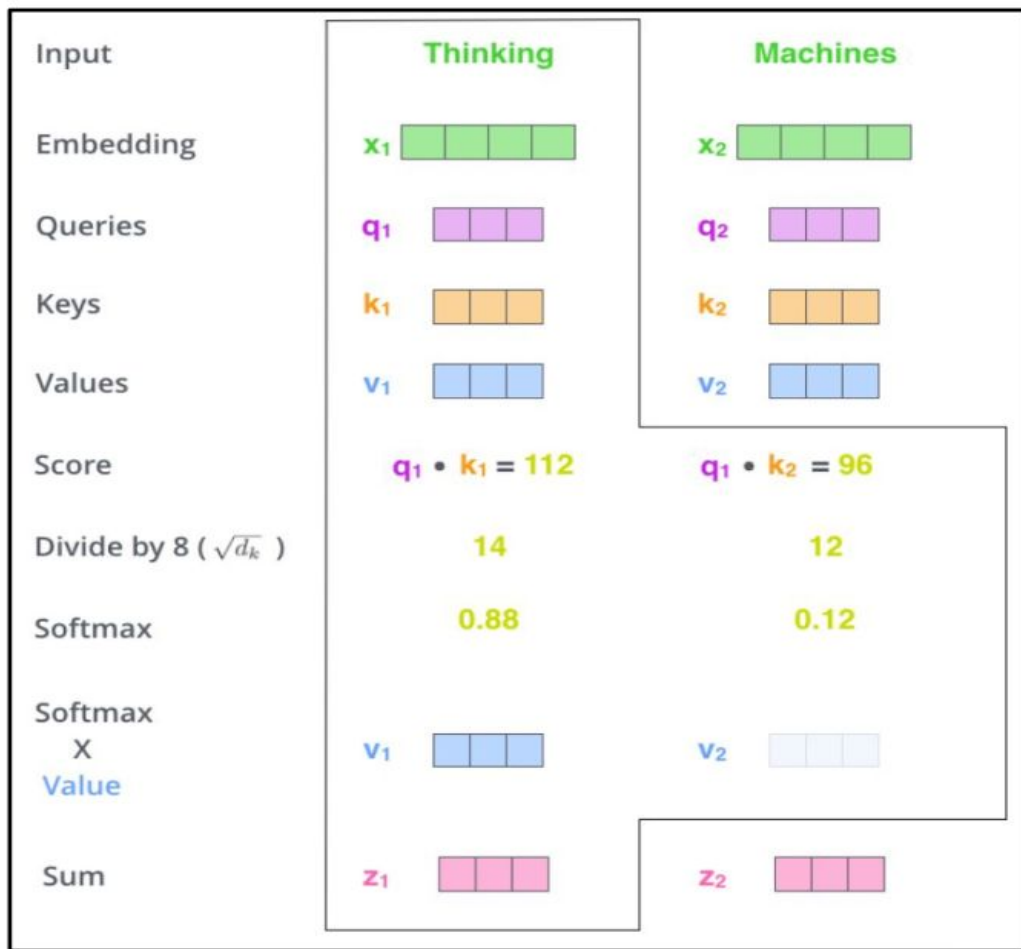
There are **8 Attention heads** ($\mathbf{Z0} \dots \mathbf{Z7} = \mathbf{Zi}$)



Backup

Self Attention

1. Find q , k and v
2. Find $\text{score}_i = q_i * k_j$
3. Normalise
4. Softmax
5. Multiply softmax output with v
6. Sum weighted value vector



Backup

Experiment – evaluation metrics

- **BLEU** (Bilingual Evaluation Understudy)
 - Evaluation metrics on how machine translation(MT) and ground truth(GT) translation are similar.

$$BLEU = BP_{BLEU} \cdot \exp \left(\sum_{n=1}^N \frac{1}{N} \log p_n \right)$$

Usually, $N=4$.

- BP_{BLEU} : penalty if multiplied when $\text{len}(\text{MT}) < \text{len}(\text{GT})$
- $p_n = \frac{\sum_i \text{the number of } n\text{-grams same in } MT_i \text{ and } GT_i}{\sum_i \text{the number of } n\text{-grams in } MT_i}$

Backup

Data Flow in Attention (Multi-Head)

1) This is our input sentence*


2) We embed each word*

3) Split into 8 heads. We multiply X or R with weight matrices

4) Calculate attention using the resulting $Q/K/V$ matrices







5) Concatenate the resulting Z matrices, then multiply with weight matrix W^O to produce the output of the layer

Thinking Machines





Input


* In all encoders other than #0, we don't need embedding. We start directly with the output of the encoder right below this one



...



W^O



Output

- **Multihead (8 headed) will increase representational power(multiple Z_i at a time)** as single attention head(Z_1) could focus only on first word.
- **Concatenate all attention heads ($Z_0 \dots Z_7$) = $Z * W^O$ = Final Z**